

Learning to Pivot with Adversarial Networks #105

Gilles Louppe, Michael Kagan and Kyle Cranmer

✉ g.louppe@uliege.be, makagan@slac.stanford.edu, kyle.cranmer@nyu.edu

🌐 <https://github.com/gloppe/paper-learning-to-pivot/>



Abstract

Several techniques for domain adaptation have been proposed to account for differences in the distribution of the data used for training and testing. The majority of this work focuses on a binary domain label. Similar problems occur in a scientific context where there may be a continuous family of plausible data generation processes associated to the presence of systematic uncertainties. Robust inference is possible if it is based on a **pivot** – a quantity whose distribution does not depend on the unknown values of the nuisance parameters that parametrize this family of data generation processes. In this work, we introduce and derive theoretical results for a **training procedure based on adversarial networks for enforcing the pivotal property (or, equivalently, fairness with respect to continuous attributes) on a predictive model**. The method includes a hyperparameter to control the trade-off between accuracy and robustness. We demonstrate the effectiveness of this approach with a toy example and examples from particle physics.

Problem statement

Data generating process $p(X, Y, Z)$

- $x \in \mathcal{X}$ are the data
- $y \in \mathcal{Y}$ are the target labels
- $z \in \mathcal{Z}$ are continuous or categorical domain labels, nuisance parameters or attribute we wish to be insensitive to.

Goal : Learn a classification/regression function f with parameters θ_f and minimizing a loss $\mathcal{L}_f(\theta_f)$ (e.g., the cross-entropy) such that

$$p(f(X; \theta_f) = s|z) = p(f(X; \theta_f) = s|z').$$

This implies that $f(X; \theta_f)$ and Z are **independent variables**.

Method

Let $r := p_{\theta_r}(z|f(X; \theta_f) = s)$ be an **adversary network** modeling the conditional distribution $Z|f(X; \theta_f)$, with parameters θ_r and loss $\mathcal{L}_r(\theta_f, \theta_r)$. As in GANs, f and r are trained simultaneously, by considering the value function

$$E(\theta_f, \theta_r) = \mathcal{L}_f(\theta_f) - \mathcal{L}_r(\theta_f, \theta_r) \quad (1)$$

that we optimize by finding the minimax solution

$$\hat{\theta}_f, \hat{\theta}_r = \arg \min_{\theta_f} \max_{\theta_r} E(\theta_f, \theta_r). \quad (2)$$




Figure 1. Architecture for the adversarial training of a binary classifier f against a nuisance parameters Z . The adversary r models the distribution $p(z|f(X; \theta_f) = s)$ of the nuisance parameters as observed only through the output $f(X; \theta_f)$ of the classifier. By maximizing the antagonistic objective $\mathcal{L}_r(\theta_f, \theta_r)$, the classifier f forces $p(z|f(X; \theta_f) = s)$ towards the prior $p(z)$, which happens when $f(X; \theta_f)$ is independent of the nuisance parameter Z and therefore pivotal.

Theoretical motivation

Assume that \mathcal{L}_f and \mathcal{L}_r are respectively set to the expected value of the negative log-likelihood of $Y|X$ under f and of $Z|f(X; \theta_f)$ under r :

$$\mathcal{L}_f(\theta_f) = \mathbb{E}_{x \sim X} \mathbb{E}_{y \sim Y|x} [-\log p_{\theta_f}(y|x)] \quad (3)$$

$$\mathcal{L}_r(\theta_f, \theta_r) = \mathbb{E}_{s \sim f(X; \theta_f)} \mathbb{E}_{z \sim Z|s} [-\log p_{\theta_r}(z|s)] \quad (4)$$

Proposition 1. If there exists a minimax solution $(\hat{\theta}_f, \hat{\theta}_r)$ for Eqn. (2) such that $E(\hat{\theta}_f, \hat{\theta}_r) = H(Y|X) - H(Z)$, then $f(\cdot; \hat{\theta}_f)$ is both an optimal classifier and a pivotal quantity.

Proof. (sketch)

- (i) For fixed θ_f , adversary r is optimal at $\hat{\theta}_r = \arg \max_{\theta_r} E(\theta_f, \theta_r) = \arg \min_{\theta_r} \mathcal{L}_r(\theta_f, \theta_r)$, in which case $p_{\hat{\theta}_r}(z|f(X; \theta_f) = s) = p(z|f(X; \theta_f) = s) \forall z, s$, and $\mathcal{L}_r \rightarrow \mathbb{E}_{s \sim f(X; \theta_f)} [H(Z|f(X; \theta_f) = s)]$;
 - (ii) Value function E can be restated as a function of θ_f only : $E'(\theta_f) = \mathcal{L}_f(\theta_f) - H(Z|f(X; \theta_f))$;
 - (iii) We have the lower bound $H(Y|X) - H(Z) \leq \mathcal{L}_f(\theta_f) - H(Z|f(X; \theta_f))$ where the equality holds at $\hat{\theta}_f = \arg \min_{\theta_f} E'(\theta_f)$ when :
 - $\hat{\theta}_f$ minimizes the negative log-likelihood of $Y|X$ under f , which happens when $\hat{\theta}_f$ are the parameters of an optimal classifier. Then, \mathcal{L}_f reduces to minimum value $H(Y|X)$.
 - $\hat{\theta}_f$ maximizes the conditional entropy $H(Z|f(X; \theta_f))$, since $H(Z|f(X; \theta_f)) \leq H(Z)$.
- \Rightarrow By assumption, the bound is active, hence $H(Z|f(X; \theta_f)) = H(Z)$ because of the second condition, which happens when Z and $f(X; \theta_f)$ are independent variables. \square

In practice

The assumption of existence of an optimal and pivotal classifier may not hold because Z directly shapes the decision boundary. However, the value function E can be rewritten as

$$E_\lambda(\theta_f, \theta_r) = \mathcal{L}_f(\theta_f) - \lambda \mathcal{L}_r(\theta_f, \theta_r) \quad (5)$$

where $\lambda \geq 0$ is a hyper-parameter controlling the trade-off between the performance of f and its independence with respect to the nuisance parameter.

- Setting λ to a large value preferably enforces f to be pivotal.
 - Setting λ close to 0 constrains f to be optimal.
- The optimal choice of λ should be **guided by a higher-level objective**.

Toy example

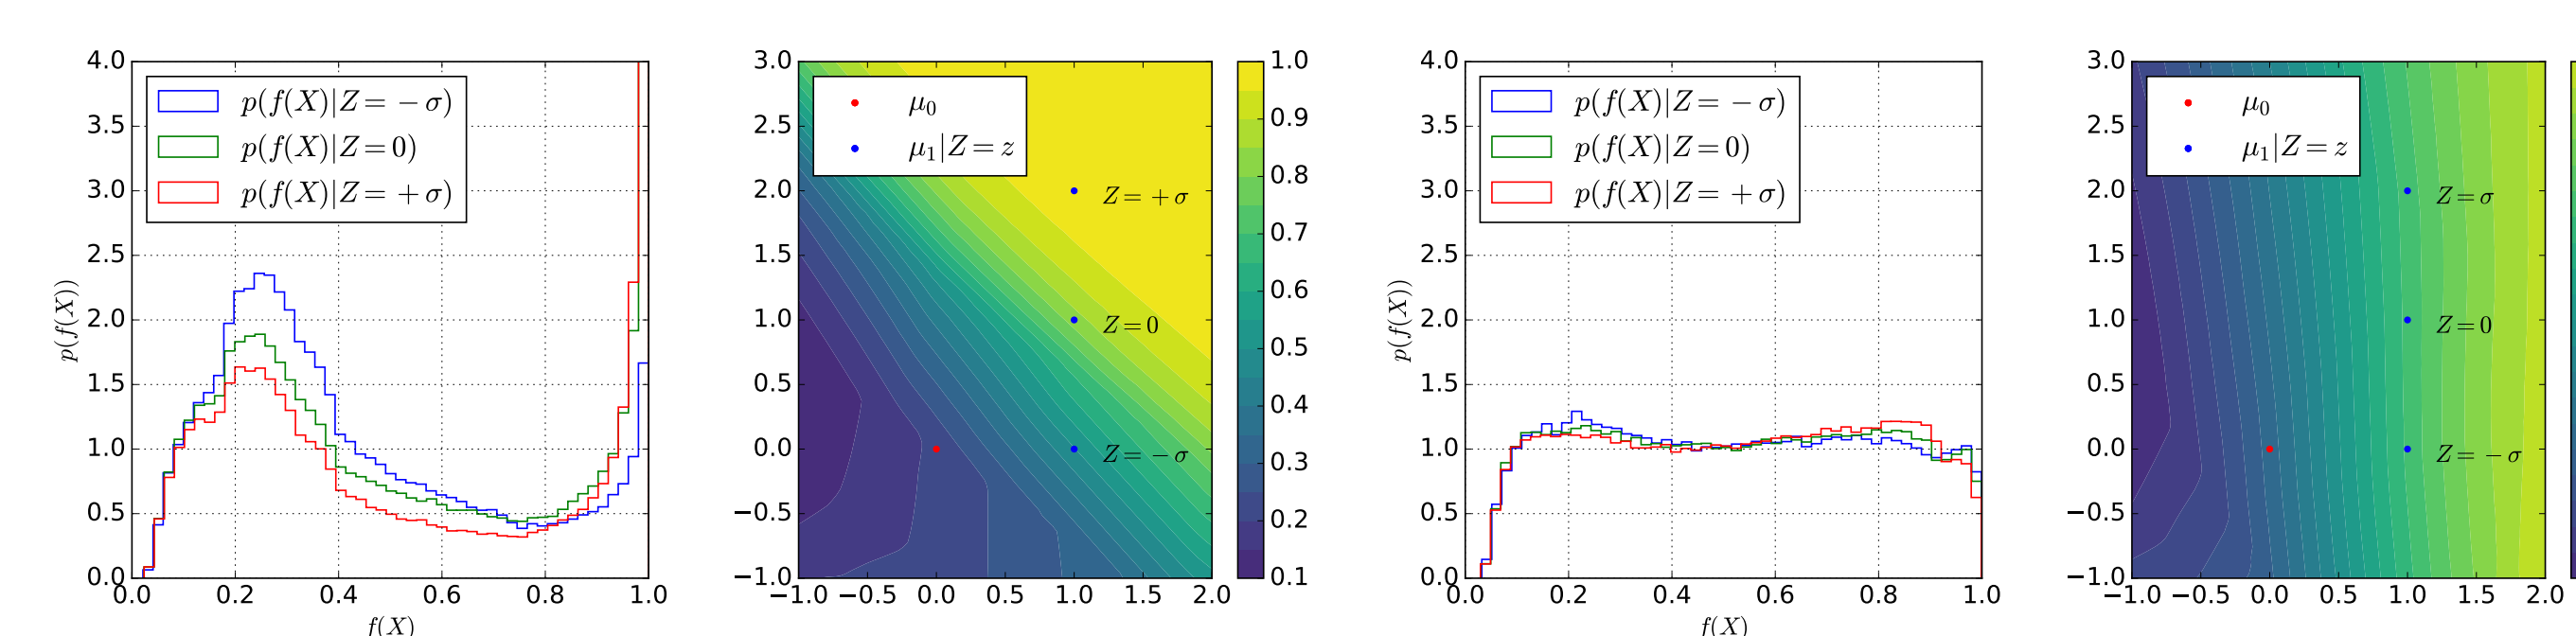


Figure 2. (Left) Decision scores without adversarial training, showing clear dependence on Z (Middle left) Decision surface. Samples are easier to classify for values of Z above σ . (Middle right) Decision scores with adversarial training. The resulting densities are now almost identical to each other, indicating only a residual dependency on Z . (Right) Adversarial training reshapes the decision function vertically to erase the dependency on Z .

Examples




Figure 3. Training curves for the toy example. Initialized with a pre-trained classifier f , adversarial training was for 200 iterations, mini-batches of size $M = 128$, $K = 500$ and $\lambda = 50$




Figure 4. Physics example. Approximate median significance as a function of the threshold on f . At $\lambda = 10$, **trading accuracy for independence results in a net benefit in terms of statistical significance**.


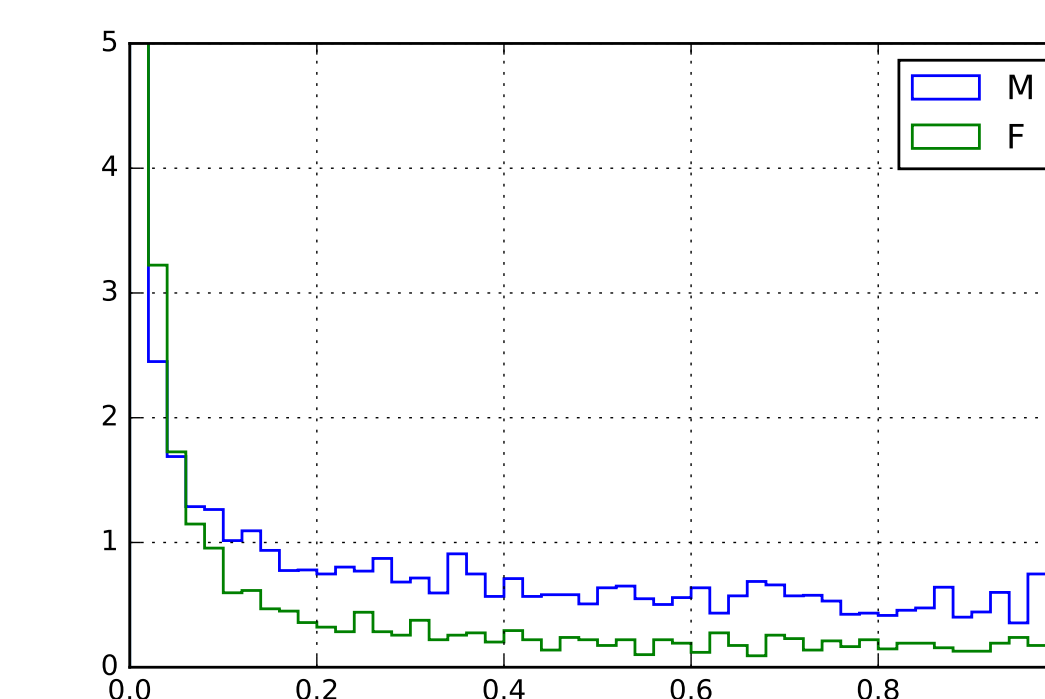


Figure 5. Building a fair classifier independent of gender. (Left) Without adversarial training, decision scores depend on gender. (Right) With adversarial training.

Conclusions

- ✓ Proposed a flexible learning procedure **for building a predictive model that is independent of continuous or categorical nuisance parameters** by jointly training two neural networks in an adversarial fashion.
- ✓ Motivated the algorithm by showing that the minimax solution corresponds to a predictive model that is both optimal and pivotal (if that models exists) or for which one can tune the trade-off between power and robustness.
- ✓ Can be used in any situation where the training data may not be representative of the real data the predictive model will be applied to in practice.
- ✓ Enable one to train predictive models that are insensitive to systematic uncertainties parametrized by continuous or categorical nuisance parameters.
- ✓ Extend to cases where independence with respect to observed random variables is desired, **as in fairness or scientific analysis**. E.g., this has been used in jet tagging (Shimmin et al, 2017) for decorrelating a model against a **continuous** attribute. The proposed approach is also **already used in several ongoing studies at the LHC**.

KC and GL are both supported through NSF ACI-1450310, additionally KC is supported through PHY-1505463 and PHY-1205376 and GL through the NRB Big Data Chair. MK is supported by the US Department of Energy (DOE) under grant DE-AC02-76SF00515 and by the SLAC Panofsky Fellowship.